

Complex Systems and DevOps: Deliverable 2

Date: November 25, 2024

Course Details

Course Name: Complex Systems and DevOps **Course Code:** 62582
Semester: Fall 2024

Team Members

Name	Student Number
Christoffer Fink	<i>s205449</i>
Kasper Falch Skov	<i>s205429</i>
Johan Søgaard Jørgensen	<i>s224324</i>
Henrik Lynggaard Skindhøj	<i>s205464</i>
Kevin Wang Højgaard	<i>s195166</i>
Sebastian Halfdan Lauridsen	<i>s215769</i>

Link to GitHub Project

Table of Contents

1. Introduction
 - Project Scope and Objectives
 - Problem Statement and Solution Overview
 - Methodology
- I. Analysis
 2. Domain Analysis
 - User Stories and Requirements
 - System Architecture Overview
 - Technical Stack Selection
 - Security Requirements
 3. Technical Foundation
 - Framework Selection Rationale
 - Development Environment Setup
 - Project Structure
- II. Implementation and DevOps Practices

4. Backend Development
 - Quarkus Framework Implementation
 - REST API Design with Siren Hypermedia
 - Database Integration
 - Business Logic Implementation
 - Security Implementation
 - JWT Authentication
 - Testing Strategy
 - JUnit Implementation
 - REST-assured Testing
 - OpenAPI Documentation
5. Frontend Development
 - React Application Structure
 - TypeScript Integration
 - Vite Build Tool Implementation
 - State Management
 - Component Architecture
 - Security Features
 - Token Security Implementation
 - Package Management
6. DevOps Implementation
 - Version Control Practices
 - Git Workflow
 - GitHub Integration
 - Continuous Integration/Continuous Deployment
 - GitHub Actions Configuration
 - Build Server Setup
 - Testing Pipeline
 - Containerization
 - Docker Implementation
 - Container Registry
 - Cloud Deployment
 - Google Cloud Setup
 - Netlify Frontend Deployment
 - Monitoring and Maintenance
7. Conclusion
 - Project Outcomes
 - Future Improvements
 - Lessons Learned

```
+- csdg8 +- auth | +- AuthController.java | +- AuthResource.java | +-
AuthService.java | +- dto | | +- CreateTokenRequest.java | | +- CreateTo-
kenResponse.java | | +- GetAuthResponse.java | | +- RefreshAccessToken-
Request.java | | +- RefreshAccessTokenResponse.java | +- TokenService.java
+- user | +- dto | | +- CreateUserRequest.java | | +- GetCollectionUser-
Response.java | | +- GetUserResponse.java | +- UserController.java | +-

```

User.java | +- UserResource.java | +- UserService.java | ...

*Snippet of the backend file structure showing its organization by the domains "auth" and "u

To further simplify new features or changes our general class hierarchy is as follows

1. Resource
2. Controller
3. Service
4. Business Logic

The resource classes only have direct interaction with a single controller in the same domain.

Another feature of this model is our separation of DTOs from implementation classes. None of

A further improvement to this structure would be, in the same vein as our DTOs, shielding the

Security Implementation

quarkus security with rolesallowed, 401 vs 403.

To secure our API we use Quarkus Security Jakarta Persistence. Our User entity's variables h

```
```java
@Entity
@Table(name = "app-user")
@UserDefinition
public class User extends PanacheEntity {

 @Username
 public String username;

 @Password
 public String password;

 @Roles
 public Set<String> role;

 //...
}
```

The `@UserDefinition` tells our application that this entity is a source of identity information, whilst the accompanying `@Username` indicates that this field is a username, the `@Password` indicates that this field is a hashed password and finally the `@Roles` indicates that this field is a collection (a `Set` in our case to force unique roles) of roles.

To secure a specific endpoint we use the `@RolesAllowed` annotation with a

single or multiple specified roles. In the following example users accessing the `/auth/token/refresh` endpoint must have either the “user” or “admin” role.

```
@POST
@RolesAllowed({ "user", "admin" })
@Path("/token/refresh")
@Consumes(MediaType.APPLICATION_JSON)
public Entity refreshAccessToken(RefreshAccessTokenRequest request) throws Siren4JException
 //...
}
```

*Simplified endpoint which uses the `@RolesAllowed` annotation to limit access.*

The annotation can also be used for more fine-grained control on service methods if a part of the application must be extra secure, however we have not utilized this functionality yet.

**JWT Authentication** For authentication we decided to roll our own JWT tokens via the MicroProfile JWT RBAC specification. Quarkus conveniently provides a library for this named `quarkus-smallrye-jwt`. In short, a JWT token is a server provided signed token which anyone can verify was signed by the server, meaning the contents and access which it grants are valid. To accomplish this JWT tokens are signed with a *private key* by the server and can be verified with the linked *public key*. Contrary to a regular encrypted transaction where the *public key* signs some data which can then only be unlocked with the *private key*. In our self-rolled JWT implementation, when a registered user sends their username and password in a POST request to the `/auth/token` endpoint, the credentials are validated and a JWT token is generated via the *private key*.

```
public String generateAccessToken(User user) {
 return Jwt.issuer(this.issuer)
 .upn(user.getUsername())
 .subject(user.id.toString())
 .groups(user.getRole())
 .expiresIn(Duration.ofMinutes(5))
 .sign();
}
```

*Our self-rolled JWT token generation.*

Of note here is the `upn` which is our main unique ID in our JWT tokens, but we also have a `subject` claim which we use to identify a user in our backend.

Once the JWT token is generated and signed it is returned to the user. When the user then sends a request to a locked endpoint (recall the `@RolesAllowed` annotation) the token is automatically verified by the Quarkus framework with the *public key* and verified if it was actually signed by the server. If it is valid and if the user has the required roles the request may proceed.

As mentioned earlier we utilize a access-token (JWT) and refresh-token pair. With the access-token being short-lived and the refresh-token being long lived. The refresh-token is a simple generated UUID and the backend keeps track of which UUIDs it generates. A future improvement will be storing the list of these generated UUIDs in persistent storage instead of in-memory as it is now. If the server terminates for whatever reason the list of valid refresh-tokens is lost and all users must re-login to get a new refresh-token. The access-token should not be persisted as it is assumed, because it is cryptographically signed, that it is always valid if it is not expired. To re-emphasize, a JWT token can only come from the server as it is the only one with the *private key*.

### Testing Strategy

The testing strategy verifies the functionality of various classes such as the Resource classes, Controller classes, and Service classes. Each class type has defined objectives to make sure the system works the correct way, which includes ensuring resources accessibility from the root.

**Resource Classes** Tests for Resource classes validate compliance with hypermedia principles and the functionality of core endpoints. These include verifying accessibility from the root, validating hypermedia using `SirenAssertion.java`, and ensuring the correct operation of core endpoint flows.

**Controller Classes** The tests for the controller classes validates data transfer objects or (DTOs) using Jakarta validation annotations. These test will enforce constraints on incoming data and rejecting invalid inputs with the appropriate error responses. The `CoinFlipControllerTest.java` gives an example to this approach by validating input data before invoking the business logic.

**Service Classes** The service classes validate the business logic. This includes verification of game mechanics, such as the coinflip logic. Thereby ensuring accurate token generation and authentication handling. The tests will address the edge cases and confirm the expected outcome.

**Strategy Evolution** The testing strategy was formalized during the project and is consistently applied to new code. Earlier parts of the project may deviate due to the absence of a defined strategy at the time. The `CoinFlip` test classes is the most modern thereby reflecting our testing strategy as the most recent iteration.

**OpenAPI documentation** Our backend uses OpenAPI documentation to provide a clear, standardized description of all exposed REST API endpoints. This documentation is automatically generated and served using the `quarkus-smallrye-openapi` library, which integrates seamlessly with our Quarkus framework. By leveraging OpenAPI, we ensure that both internal and external de-

velopers have access to a comprehensive, up-to-date overview of our API. All API endpoints, including their HTTP methods (GET, POST, etc.), paths, request parameters, and response formats, are automatically documented based on annotations in the codebase. For example:

```
@Path("/game/coin-flip")
@Produces(Siren4J.JSON_MEDIATYPE)
public class CoinFlipResource {

 @POST
 @RolesAllowed("user")
 public Response play(PlayCoinFlipRequest request) {
 return this.coinFlipController.play(request);
 }
}
```

This defines a POST endpoint, `/game/coin-flip`, which consumes and produces JSON. Quarkus uses this metadata to populate the OpenAPI documentation.

The OpenAPI specification is made available at the endpoint:

- <http://localhost:8080/q/openapi>

An interactive Swagger UI is also available at:

- <http://localhost:8080/q/swagger-ui>

By integrating OpenAPI documentation into our backend, we have established a robust and developer-friendly way to manage, explore, and share our API. This not only enhances collaboration but also ensures our API remains well-documented and easy to consume. The use of OpenAPI further supports our project's DevOps practices, enabling automation in testing and deployment workflows.### 5. Frontend Development ##### React Application Structure

The project is a frontend application built using React with TypeScript, designed to provide an engaging and interactive user experience. It features a landing page with multiple sections, including clear navigation to different game pages. The application incorporates secure authentication with login and registration pages. Each game is implemented on its own dedicated page, leveraging a modular and structured approach. This design ensures scalability, allowing for seamless project expansion or simplification, while maintaining clean and maintainable code.

**Key sections and features** The landing page is the first page of the application and where the user will start when they enter. It is built up by using several components that form the landing page when put together.

```
function LandingPage() {
 return (
```

```

 <div className="min-h-screen flex flex-col">
 <Header />
 <HeroSection />
 <GamesSection />
 <StartSection />
 <Footer />
 </div>
);
}

```

```
export default LandingPage;
```

The **header** is at the top of the landing page, and serves as a navigation bar. It has the logo of the application and then links to different sections of the application: **login**, **register** etc. Below that is the **HeroSection**, where you would typically have a banner, welcoming the user onto the page and inviting them to navigate around the application to explore the different features. Further down is the **GamesSection** where the different implemented games are listed along with buttons that navigate to said games. Below that there is a “Start” section, which is the component that invites the user to sign up in order to play the games. Lastly a **Footer** is implemented, which is a generic footer containing copyright information and links to resources such as terms of service and privacy policy. The application also has separate game pages, which when navigated to, shows each of the games.

**Navigation and Routing** For routing in the application **react-router-dom** is used to enable navigation between different pages. The routes are defined in the **index.tsx** file, and currently contains routes for **blackjack**, **coinflip**, **register** and the landing page.

```

ReactDOM.createRoot(document.getElementById('root') as HTMLElement).render(
 <React.StrictMode>
 <Router>
 <AuthProvider>
 <UserProvider>
 <Routes>
 <Route path="/" element={<LandingPage />} />
 <Route path="/coinflip" element={<CoinFlipPage />} />
 <Route path="/poker" element={<PokerPage />} />
 <Route path="/blackjack" element={<BlackjackPage />} />
 <Route path="/register" element={<RegisterPage />} />
 </Routes>
 </UserProvider>
 </AuthProvider>
 </Router>
 </React.StrictMode>
)

```

);

This setup makes it easy if new games are added, because they can just be added to the index file.

**TypeScript Integration** The project is built with TypeScript, which makes it possible to do static type checking. This is super helpful in order to catch errors during development instead of having to catch them at runtime. TypeScript and the static type checking is best used when defining interfaces and types for components props and states. This limits the data passed between components such that it has to be structured properly for the component to accept it.

**Vite Build Tool Implementation** For building the application this project uses Vite. This is a popular build tool with some advantages mainly aimed at ease of use for the developers. The advantages include a really fast development server providing instant updates when changes are made to the code. It supports TypeScript out of the box, which minimizes the amount of work developers have to focus on configuring build pipelines.

**State Management** In this project, React's built in state management hooks are used. UseNavigation from react-router-dom is used to handle the navigation state and allows the user to navigate to different pages. Local states are also being used within the components with useState. useState can be used to handle specific UI states, for example when users are interacting with the application, or when dynamic content rendering is used.

In our application, we manage state efficiently using React's useContext in combination with a Provider. This approach centralizes state management, making the application easier to scale, maintain, and extend as the complexity of features grows. With this setup: \* State stores the current application data. \* Dispatch handles updates and actions, ensuring predictable and controlled state transitions.

By wrapping our components in context Providers, any part of the application can access and modify the shared state without the need for cumbersome prop drilling, enhancing the overall developer experience. We have implemented specific contexts and providers tailored to key areas of our application: \* Authentication Context: Manages authentication tokens, login status, and user session data. \* User Context: Handles user-specific information, such as profile details or preferences. \* Blackjack Context: Manages the state for our Blackjack game, including game logic, player actions, and dealer interactions.

This modular approach ensures that each context is focused on its specific domain, improving code organization and maintainability. By leveraging useContext and the Provider, we create a robust and scalable foundation for managing



complex state transitions and interactions across the application while adhering to React's functional component paradigm.

## Security Features

**Token Security Implementation** In our frontend application, we have chosen to store authentication tokens securely in cookies. This decision strikes a balance between usability and security, ensuring sensitive token data is protected while providing a seamless user experience. The main reason cookies are optimal for token storage is their ability to include security features like the `HttpOnly` flag. This flag makes cookies inaccessible to JavaScript, significantly reducing the risk of XSS (Cross-Site Scripting) attacks. Additionally, the `Secure` flag ensures cookies are transmitted only over HTTPS, preventing exposure over unencrypted connections. For token expiration, we enforce short lifetimes to limit the risk of misuse if a cookie is compromised. While the expiration mechanism is managed on the backend, it adds another layer of security by ensuring tokens do not remain valid indefinitely. We store both the refresh token and access token in cookies. The frontend includes functionality to refresh the access token as its expiration approaches, ensuring uninterrupted user sessions. This mechanism allows users to remain logged in as long as they keep the site open. Moreover, cookies enable persistent login between visits. If the user closes the webpage and later reopens it, the application automatically checks if valid tokens exist in the cookies. If the tokens are not expired and can still be used, the user is logged back into the site seamlessly. By leveraging cookies for token storage and refresh management, we provide a secure and user-friendly authentication experience. ##### **Input Validation and Sanitation** We perform input validation and sanitization in the frontend to protect our backend from malicious inputs and to catch invalid data early in the process. This proactive approach enhances security while improving the user experience by providing immediate feedback if the input does not meet our predefined standards. For instance, we enforce specific requirements for fields such as passwords and usernames, ensuring they comply with our security and usability guidelines before they are submitted to the server. To clarify, we also sanitize and validate inputs in the backend, but catching validation errors at the frontend enhances the overall system by reducing invalid requests sent to the server. ##### **Secure Communication (HTTPS)** Our frontend exclusively uses HTTPS to ensure that all data transmitted between the client and the server is encrypted. This protects sensitive information, such as authentication tokens and user data, from interception or tampering during transmission, maintaining the integrity and confidentiality of the communication.

**Package Management** In our project, we utilize npm (Node Package Manager) as the primary package manager for frontend development. npm is a powerful tool that simplifies the process of managing dependencies and libraries required for the project. By specifying the necessary packages in the `package.json`

file, npm allows us to efficiently install, update, and manage all dependencies in a structured and repeatable manner.

One of the key advantages of npm is its ability to define and execute custom scripts. These scripts streamline various development and build tasks. For example:

npm run dev: This command is configured to start a development server, providing a live-reloading environment that makes it easier to test and iterate on frontend code during development. npm run build: This command specifies the steps required to build the production-ready version of the frontend application. It typically compiles, minifies, and optimizes the code for deployment. Additionally, npm enhances the security of our project by providing the npm audit tool. This tool scans the project dependencies for known vulnerabilities and security threats. By running npm audit, we can:

- Detect security issues in real-time across all installed libraries.
- Receive detailed reports on the nature of vulnerabilities, their severity, and potential fixes.
- Apply automated patches for certain vulnerabilities using npm audit fix.

By regularly auditing our dependencies, we ensure that our project remains secure and compliant with best practices, reducing the risk of threats from third-party libraries. As a future improvement, we could set up a GitHub Action to automatically create issues whenever vulnerabilities are detected during the auditing process. This automation would help us maintain a proactive approach to dependency management, ensuring timely resolution of potential security concerns.

In addition to managing dependencies and automating tasks, npm also enables us to run custom test suites. For instance, by specifying test commands in package.json, such as npm run test, we can execute different testing frameworks or tools to ensure the quality and stability of our code. This feature presents an opportunity for future improvement, as we plan to leverage it when implementing a comprehensive testing strategy for the frontend.

By leveraging npm for package management, task automation, and security auditing, we ensure that our frontend development workflow is consistent, efficient, and secure. It simplifies collaboration across the team, as all required dependencies and commands are easily accessible, while safeguarding the project against potential vulnerabilities. ##### Framework Selection Rationale

Our project uses **React** with **TypeScript** as the primary framework and language for the frontend, paired with **Vite** as the build tool. This choice is informed by the features we have discussed previously and is further supported by the following specific points:

#### **React:**

- Component-based architecture: Simplifies the development of reusable,

modular, and scalable UI elements, enabling consistent design and functionality across the application.

- **Robust ecosystem:** With a large developer community, React provides abundant resources, libraries, and third-party integrations to address common challenges efficiently.
- **Extensibility:** React integrates seamlessly with tools such as react-router-dom for routing and state management libraries like Redux or Context API, making it versatile for various project needs.

#### **TypeScript:**

- **Static type-checking:** Reduces runtime errors by catching issues during development, ensuring more robust code.
- **Enforces stricter coding standards:** Enables explicit definitions for props, states, and API responses, fostering better collaboration and reducing ambiguity in the codebase.
- **Enhanced maintainability:** Improves code readability, simplifies debugging, and makes the project easier to scale.

#### **Vite:**

- **Lightning-fast development server:** Delivers near-instant hot module replacement (HMR), significantly boosting development speed and productivity.
- **Ease of configuration:** Requires minimal setup, letting developers focus on building the application rather than configuring the tooling.
- **Optimized builds:** Produces lightweight, high-performance builds that outperform traditional tools like Webpack, ensuring a smoother user experience.

This combination of React, TypeScript, and Vite provides a well-rounded stack that prioritizes development efficiency, scalability, and code quality. It is tailored to meet the demands of a modern, interactive frontend application, making it an ideal choice for our project.

**Project Structure** Our project structure is meticulously designed to adhere to React best practices, ensuring maintainability, scalability, and readability of the codebase. We follow a component-based architecture, where each UI element is encapsulated in its own reusable and self-contained component. This approach promotes reusability and makes it easier to manage changes or updates.

We also emphasize clear separation of concerns by organizing files logically. For example, each component has its own directory containing its Typescript file, CSS (or module CSS for scoped styling), and any relevant assets. This structure ensures that related code is grouped together, making the project easier to navigate.

Another key best practice we follow is state management. For local state, we use React's `useState` and `useReducer` hooks, while global state is managed with

Context API or third-party libraries, depending on the project requirements. For instance, in our authentication flow, we use Context API to share user and authentication states across the app, avoiding unnecessary prop drilling.

Finally, we focus on writing clean, readable code by adhering to consistent formatting rules enforced by tools like Prettier and ESLint. This ensures that all developers contribute code that follows a unified style, reducing friction during code reviews.

### GitHub Workflow

- **ci.yml:** This file defines the GitHub workflow for automating key stages of our development pipeline, including building and testing the project before deployment. The workflow is triggered by specific events, such as pushing to the main branch or opening a pull request, ensuring that our codebase remains stable and free of regressions. By automating these processes, we reduce manual effort, maintain high code quality, and accelerate development cycles.

Further details about the setup, configuration, and implementation of the ci.yml file, as well as how it integrates with other DevOps practices, are discussed in the DevOps section of this report. ##### Source Directory (**src**) Our project's directory structure is carefully designed to promote clarity, maintainability, and scalability, adhering to established best practices in React development - **Components:** This folder contains individual React components, each in its own folder with associated files like .tsx. By isolating components, we ensure reusability, making it easier to share UI elements across pages. This modular approach also simplifies debugging and updating specific parts of the UI. - **Pages:** Complete pages are created by assembling multiple components. This separation of "components" and "pages" enforces a clean hierarchy, where pages focus on layout and orchestration while components handle granular functionality. This approach improves readability and helps maintain a clear distinction between reusable pieces of UI and higher-level structures. - **State:** The State folder centralizes state management, using Context API. This setup avoids scattering state-related code across the application, making it easier to debug, scale, and extend the application's logic. - **Styling:** All .css files are grouped under this directory to maintain consistency in styling and enforce a separation of concerns. Using a dedicated folder allows developers to quickly locate and modify styles, whether they are global or scoped to specific components. - **Tests:** Housing all test files in a single directory ensures that testing remains a first-class citizen in the project. It helps maintain an organized structure, where tests are easy to locate and run, and it encourages consistent testing practices across the team.

### 6. DevOps Implementation

### Version Control Practices

**Git Workflow** Git is an open-source version control system that can be used between developers to help manage and track changes in their repositories. It has become the standard tool that developers use to manage software development projects ranging from small personal projects to huge projects in big organizations. The tool helps track changes to files over time, allows several developers to collaborate on the same project while ensuring that they don't accidentally overwrite each other's code, and helps provide a history of changes which developers can also use to revert to previous iterations of the same file if problems arise. Git is also being used in this project. Currently the project consists of three repositories. One for the frontend, backend and the report. In each of these repositories each group member has their own branch in which they can make changes without interfering with the other group members work. When a group member has then fully developed or written something, it can be merged into the dev branch.

**GitHub Integration** Github is an extension to the git version control system, that also serves as a hosting platform for git repositories. It comes with a lot of features that can aid the collaboration between developers to a great extent. For this project, a github organization has been made. That way the group can have their three repositories in a shared space. To each repository the group also has a work board, where tasks can be created such that each group member can see what needs to be done. This helps create an overview of the project status. When group members have developed or written something, they can create a pull request, "asking" for their branch to be merged into a different branch. It is essentially an invitation for the rest of the group to review what has been committed to the repository before merging it into the development branch. This is a way for developers to collaborate and catch errors before they reach a shared branch. Pull requests can also be used to setup continuous integration and continuous deployment.

## **Continuous Integration/Continuous Deployment**

### **GitHub Actions Configuration**

#### **Build Server Setup**

**Testing Pipeline** For the backend of the project, we're using a GitHub Actions workflow as a testing pipeline. The workflow is triggered on push events as well as pull request events into the dev and main branch. The **build** part of the workflow has steps included to set up the JDK, generating fake JWT keys and building the project using Maven. Following that there is a **test** job which depends on the build job. This means that the test job will not run if the build job fails. If the build job has succeeded, it will download the build artifacts, set the JWT key environment variables, and run the tests for the backend using Maven. This pipeline runs, as said previously, every time a pull request is made,

making it super useful, as it prevents any pull requests that do not either build or pass all the tests from being merged into the dev or main branch. For the frontend we have a similar pipeline that first installs the dependencies using **npm run ci**, builds the solution using **npm run build**, and runs the tests using **npm run test**.

**Containerization** Containerization is a lightweight form of virtualization that packages an application and its dependencies into a single, portable unit called a container. This approach allows our applications to run consistently across various environments, from development to production, without being affected by differences in underlying infrastructure.

**Docker Implementation** Docker is used to host and build the final application that Quarkus builds, both for local development and cloud deployment. This approach ensures that the production environment closely resembles the development environment, minimizing discrepancies and potential issues.

By using Docker, we can create a consistent and isolated environment for our application, which simplifies the deployment process and enhances reliability across different stages of the software's lifecycle.

**Container Registry** For hosting the Docker container in the cloud, we utilize Google Cloud's Artifact Registry. This service provides a centralized location for storing and managing build artifacts and dependencies, which is crucial for maintaining a streamlined and integrated development workflow.

Given that our final deployment is hosted on Google Cloud Run, using Artifact Registry aligns perfectly with our infrastructure, ensuring seamless integration and efficient management of our container images.

**Cloud Deployment** Cloud deployment involves hosting applications and services on cloud infrastructure, offering scalable, flexible, and efficient resource management. By leveraging cloud platforms, we can deploy our applications across global data centers, ensuring high availability and performance. This approach reduces the need for on-premises hardware, enabling us to focus on development and innovation rather than managing infrastructures.

By deploying applications closer to users through global data centers, we can reduce latency and enhance user experience. This global reach, combined with the reliability and security of cloud services, makes cloud deployment an ideal choice for us when looking to expand our reach in the market.

**Google Cloud Backend Deployment** As mentioned in the Container Registry chapter, the deployment of the backend is hosted on the Google Cloud Platform (GCP). Why did we choose Google Cloud over Microsoft Azure, especially since DTU provides students with an Azure subscription. Initially, we

attempted to set up our cloud infrastructure on Azure but quickly encountered authentication issues when integrating our GitHub Actions workflows.

A runner/account in Azure's IAM is required, but creating one requires access to Microsoft Entra ID, formerly known as Azure Active Directory, which students don't have. Instead of dealing with this complexity and risking losing access after further development, we decided to use Google Cloud, which offers us full control over the entire cloud environment. And it is free to use

In our cloud setup, we have the following elements: ##### Google Cloud Platform (GCP) Components - **SQL Instances**: Host our Quarkus PostgreSQL database instance. - **Artifact Registry**: Store our Docker containers. - **Cloud Run**: Deploy our containerized application and expose it to the web. - **Identity and Access Management (IAM)**: Manage principals and service accounts with the proper access levels, ensuring adherence to best practices for cybersecurity.

GitHub (GH) Components

- **Security - Secrets and Variables**: Store API keys and credentials securely using Repository Secrets. We have chosen to store our secrets in GitHub to separate our configuration from our deployment.

In summary, choosing Google Cloud Platform for our backend deployment has provided us with greater control and integration capabilities, overcoming the limitations we faced with Azure. This setup ensures a secure and scalable environment, supporting our development and deployment needs effectively. The pay-as-you-go model of GCP services eliminates the need for significant upfront investments, making it a cost-effective solution. Many of the services we utilize are free of charge under the chosen plans, allowing us to optimize our budget while leveraging advanced cloud capabilities.

**Netlify Frontend Deployment** In our project, we utilize Netlify as the platform for hosting and deploying our frontend application. Netlify offers an automated and seamless deployment workflow by integrating directly with our GitHub repository. Every time changes are pushed to the main branch, Netlify automatically pulls the latest code, runs the specified build command (npm run build), and deploys the application to its globally distributed Content Delivery Network (CDN).

As part of this automated process, Netlify provides real-time feedback on the deployment status. If the npm run build command succeeds, the application is immediately deployed, and the live site is updated. If the build fails, Netlify generates detailed logs, making it easy to identify and resolve any issues. After deployment, Netlify confirms whether the site was successfully deployed, ensuring full transparency in the deployment process.

In addition, for every pull request, Netlify runs a Lighthouse audit to evaluate the performance, accessibility, best practices, and SEO of the site. This feature

provides a clear score and actionable insights, helping us continuously optimize the application before merging changes into the main branch. By catching potential issues at the pull request level, we ensure that only high-quality updates make it into production.

We have also implemented GitHub Actions as part of our CI/CD pipeline. These actions complement Netlify by automating additional checks, such as running unit tests, linting, and verifying configurations before deployment. GitHub Actions provide an extra layer of validation, ensuring that the codebase adheres to predefined quality standards and that potential issues are identified early in the development cycle.

Netlify's integration with GitHub and the use of Github Actions significantly streamline our deployment process. Features like automatic builds, deployment status notifications, and Lighthouse audits on pull requests, we are able to maintain a high-quality frontend application. Its global CDN ensures that the deployed site loads quickly for users worldwide, further enhancing the user experience. By leveraging Netlify, we have a reliable, efficient, and developer-friendly solution for managing our frontend deployment pipeline.

Another of the key advantages of using Netlify is that it offers its core features for free, including automated builds, deployments, and a globally distributed CDN. This makes it an excellent choice for small teams or projects, allowing us to deploy a high-quality frontend application without incurring additional costs.

**Monitoring and Maintenance** A key principle in DevOps is the ability to act based on metrics. To do this we use Micrometer metrics, a common monitoring facade which is vendor neutral, alike SLF4J is for logging. We use the Prometheus format as Quarkus provides a convenient library for combining these, `quarkus-micrometer-registry-prometheus`. To display and view these metrics we use Dashbuilder which allows easy visualization of the metric data via a YML format. Quarkus also has minimal config library for this, named `quarkus-dashbuilder`.

We have also looked into setting up a centralized logging system, but due to time constraints were unable to implement it in time for the writing of this report.

Another clear improvement is monitoring if our frontend server is up or overloaded.

## 7. Conclusion

**Project Outcomes** Our project successfully addresses all course objectives through the implementation of a secure, scalable, and fully functional full-stack application.

We developed the frontend using React with TypeScript, leveraging its modular, component-based architecture to build reusable UI components and manage



application state. For the backend, we utilized Quarkus, a container-first Java framework, to create a REST API that handles business logic, user authentication, and persistent storage. Quarkus provided us with built-in dependency injection and Jakarta validation, streamlining development.

Application state was managed effectively on both the frontend and backend. On the frontend, we used the React Context API to maintain global state, such as user authentication and game results, without unnecessary prop drilling. On the backend, the PostgreSQL database served as the single source of truth, persisting user and game data.

To secure the backend, we implemented JWT authentication with role-based access control. Users receive short-lived access tokens and long-lived refresh tokens, ensuring secure access to protected endpoints while minimizing the need to re-enter credentials. This approach effectively supports role-based restrictions, such as limiting game endpoints to authenticated users.

In tackling the complexity of a gambling platform, we applied multiple frameworks, including Quarkus for backend development and Siren4J for implementing hypermedia-driven API responses. These tools allowed us to structure the backend effectively and dynamically guide clients through API interactions without hardcoding paths.

We used Maven to manage backend dependencies and npm for the frontend, streamlining development with tools for version management and package installations. For testing and deployment, we integrated GitHub Actions to automate CI/CD pipelines. This ensured that every code push was built, tested, and deployed automatically, reducing manual effort and maintaining high-quality code.

The backend was containerized using Docker, enabling consistent deployments across environments. The frontend was deployed on Netlify, leveraging its free tier for hosting and automated builds. These platforms facilitated a seamless deployment process and ensured scalability.

By integrating all these elements, we designed and implemented a fully functional full-stack application, meeting the course objectives and demonstrating a practical application of the knowledge gained throughout the project.

## **Future Improvements**

## **Lessons Learned**